



# Segmenting Retinal Blood Vessels with Deep Neural Networks



A decorative pattern of hexagons in various shades of blue and cyan. Some hexagons contain icons: a lightbulb, a thumbs up, a smartphone, a magnifying glass, and a gear. A network diagram with a central node and five peripheral nodes is also visible.

1

# { Problem statement }




What are trying to do? And why?



# Why?

- The condition of the vascular network of human eye is an important diagnostic factor.
- Segmentation is a nontrivial task due to variable size of vessels, low contrast, and potential presence of pathologies.
- Deep Networks are known to perform well in computer vision, but are hard to train on small datasets.
- They are usually inefficient both in terms of training and inference time.

What do we want to achieve?





# What data do we have?

## DRIVE dataset

- ◆ 20 training images
- ◆ 20 test images
- ◆ 560 x 560 pixels


## HRF dataset

- ◆ 15 images (healthy) +
- ◆ 15 images (diabetic retinopathy) +
- ◆ 15 images (glaucomatous)
- ◆ 3504 x 2336 pixels
- ◆ 6 random images chosen as a test set




# What Methods do we use?

convnets





# U-Net



A decorative pattern of hexagons in various shades of blue and cyan. Some hexagons contain icons: a lightbulb, a thumbs up, a smartphone, a magnifying glass, and a gear. A network of dots is also visible on the left side.

2

# { Preprocessing }

How to handle the data?





# Overview

## Channel-Wise normalization

Compute dataset-wide mean and std for each channel, and then use that to normalize.

## Extracting random crops

For each image, random crops of size 64x64 or 128x128 are extracted. We generate between 150 and 500 random patches per image.

## Random Rotations

Each image is cloned 3 to 6 times, and each copy is randomly rotated. Rotation angle is chosen between 0 and 270 deg.

## Patch-wise normalization

For each patch, and for each channel, subtract its mean and divide by its std

## Conversion to grayscale


Convert RGB image to grayscale.

## ZCA - whitening

Decorrelates pixels in the image.

# Generating random patches



- Essential to obtain enough data points
- Acts as data augmentation
- Combined with random rotations



x ~300

# Patch-wise normalization

- Proved to be extremely helpful
- Easy to compute
- In the end the only color transformation we used






# Patch-wise normalization (code)

```
def ChannelWiseStd(with_mean=True, with_std=True):  
    def f(images, masks):  
        for i, img in enumerate(tqdm(images, desc="{:<10}".format("Std"))):  
            mu = np.mean(img, axis=(0, 1), keepdims=True)  
            sigma = np.std(img, axis=(0, 1), keepdims=True) + 1e-6  
  
            if with_mean:  
                img = img - mu  
            if with_std:  
                img = img / (sigma+1e-6)  
  
            images[i] = img  
  
        return images, masks  
    return f
```

# ZCA whitening



- Decorrelates the features
- Normalizes variance of each feature
- Requires computing SVD -- we couldn't use it due to memory requirements :(
- Others report great results -- perhaps we should try to adopt it after all?

A series of hexagonal icons in various shades of blue and cyan are arranged along the left edge of the slide. The icons include a lightbulb, a thumbs-up, a network node, a smartphone, a magnifying glass, a gear, and a speech bubble.

3

# { Neural Network Details }

How exactly our net looks like? How is it trained?



# Overview

## **Batch-Norm**

Deals with covariate shift.  
Greatly speeds up training.

## **Lots of conv layers**

We use several different  
building blocks

## **Dropout**

Standard way to combat  
overfitting

## **Adam**

Because we love adaptive  
learning rate.

## **Up-scaling with convolution.**

Each up-scaling operation if  
followed by a  $1 \times 1$   
convolution.

## **U-Net design choices**

Lot's of those...



# ConvNets

Cool demos:

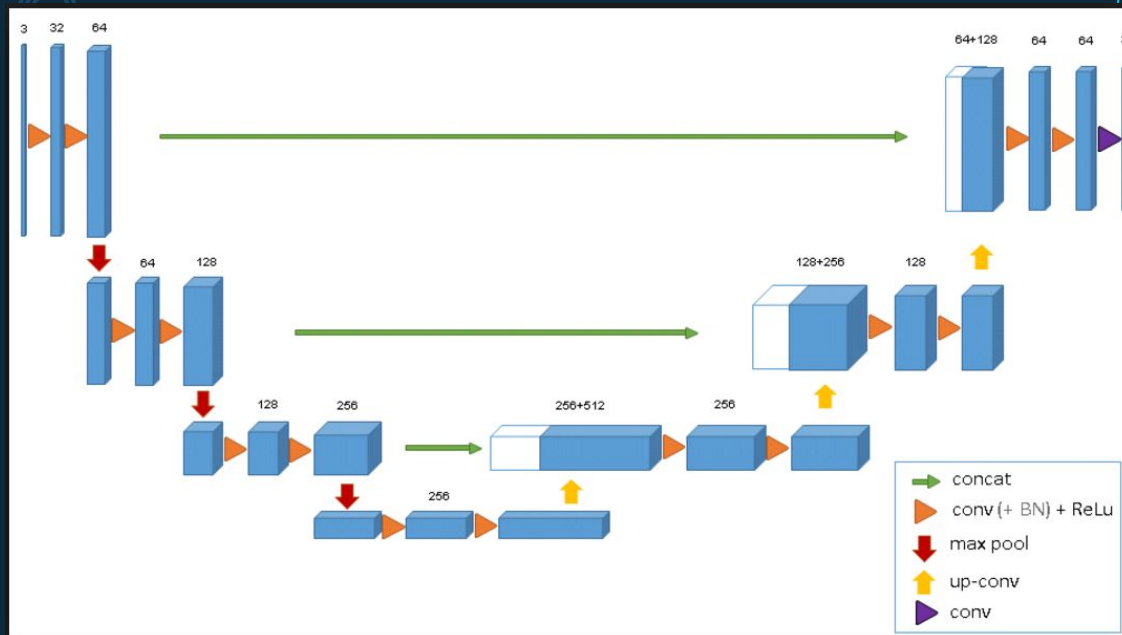
<http://cs231n.github.io/convolutional-networks/>

<http://scs.ryerson.ca/~aharley/vis/conv/>



# Design Choices

- 3 levels instead of 4: we use smaller crops
- Upsampling with  $1 \times 1$  conv to preserve the shape
- Different number of channels
- Dropout only in DoubleConv



# Adam



$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t.$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2.$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}.$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}.$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t.$$



# Batch Norm

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

A decorative pattern of hexagons in various shades of blue and cyan on the left side of the slide. Some hexagons contain icons: a lightbulb, a thumbs up, a smartphone, a magnifying glass, and a gear. A network of dots is also visible.

4

## { Results }

How does it work?



# Overview

**How did we perform?**

Can we compare to anyone?

**Prediction on fullsize images**

We train on random patches. How do we predict on full images?

**Evaluation measures**

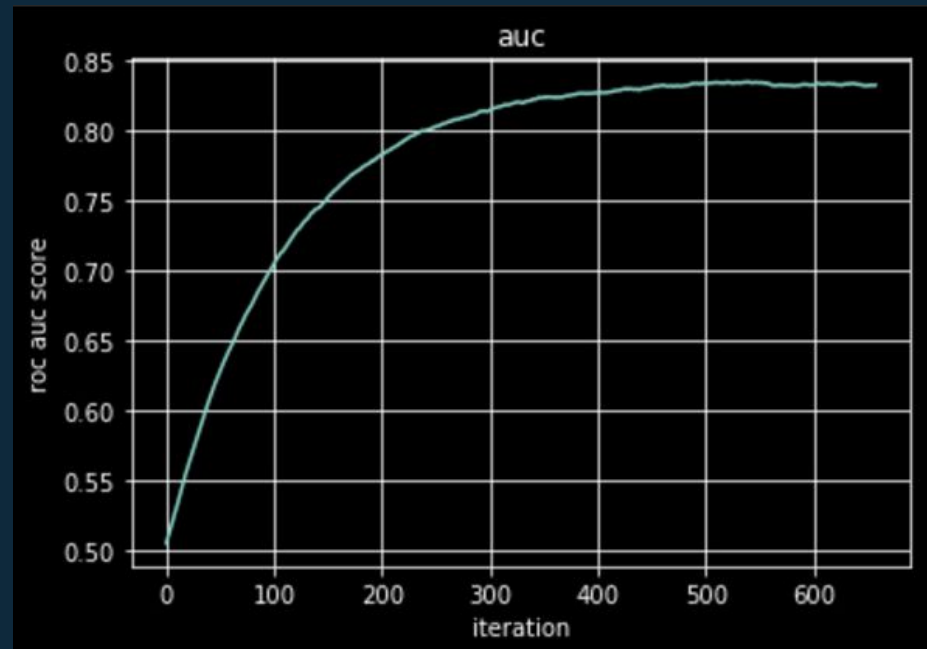
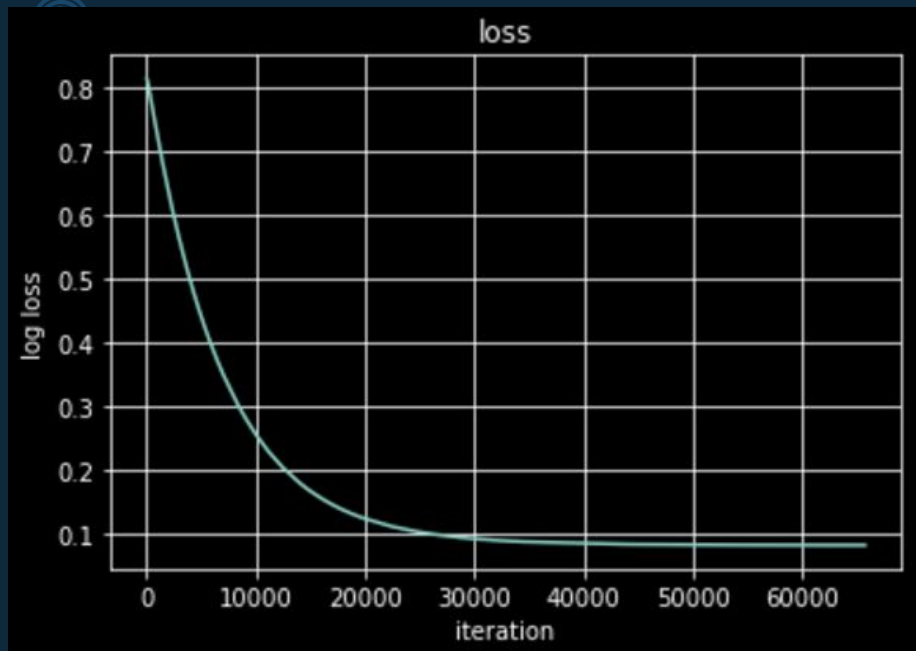
How do we measure network's performance?

**Training progress**

How to visualize learning?



Training progress (smoothed with IIR 1st order filter)





# Performance measures

## Accuracy


- Very intuitive
- Sensitive to class imbalance

## AUC (ROC)

- Robust to class imbalance
- Expensive to compute
- Measured as area under a curve, where we Plot TPR vs FPR at various thresholds


## Log Loss

- This is what we directly optimize
- Not really intuitive (KL divergence)

























# Results

LOSS \ DSET	DRIVE	HRF
Log Loss	0.075	0.084
Accuracy	0.971	0.966
AUCROC	0.911	0.828



# { Tools we used }

5



"If I have seen further, it is by  
importing from the code of  
giants."

- Definitely Not Isaac Newton



# Overview

## Visdom

Facebook's lightweight alternative to Tensorboard.

Supports plotly graphs and video exporting.

Used to track experiments

## PyTorch

New deep learning library straight from Facebook AI Research.

Numpy-like, with strong GPU acceleration.

Provides Autograd engine.

## NumFocus Stack

Jupyter + NumPy

## PIL.Image

Please, stop using OpenCV when it's not absolutely necessary!

# PYTORCH



## Back-propagation uses the dynamically built graph

```
from torch.autograd import Variable
```

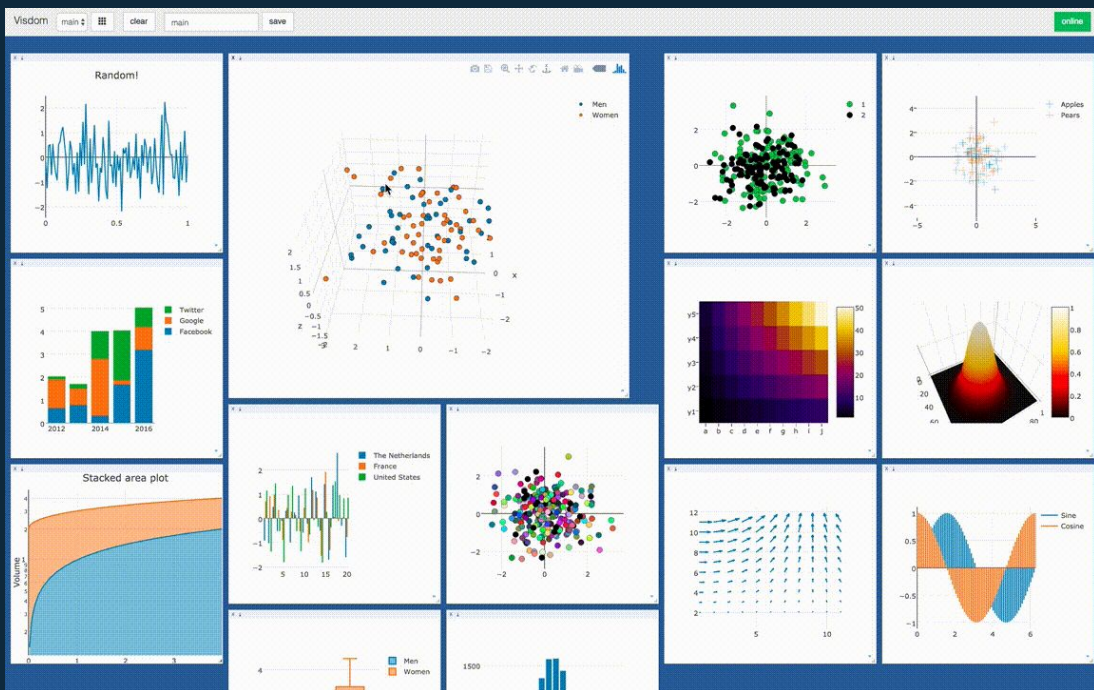
```
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))
```

```
i2h = torch.mm(W_x, x.t())  
h2h = torch.mm(W_h, prev_h.t())  
next_h = i2h + h2h  
next_h = next_h.tanh()
```

```
next_h.backward(torch.ones(1, 20))
```



# Visdom











# Pillow

- Great alternative to OpenCV
- Clean and up-to-date docs
- Many useful primitives for manipulating images
- pip installable!

🏠 Pillow (PIL Fork)

4.0.x

Installation

Handbook

☐ Reference

☐ Image Module

☐ Examples

☐ Functions

The Image Class

Attributes

[Docs](#) » [Reference](#) » Image Module

[🔗 Edit on GitHub](#)

## Image Module

The `Image` module provides a class with the same name which is used to represent a PIL image. The module also provides a number of factory functions, including functions to load images from files, and to create new images.

## Examples

The following script loads an image, rotates it 45 degrees, and displays it using an external viewer (usually xv on Unix, and the paint program on Windows).



A decorative pattern of hexagons in various shades of blue and cyan on the left side of the slide. Some hexagons contain icons: a lightbulb, a thumbs-up, a network node, a smartphone, a magnifying glass, a gear, and a speech bubble.

6

{ Simple baseline}



What have we used?



**OpenCV**






# How does it work?

1. **Input image preprocessing**
2. **Algorithm**
3. **Output image processing**
4. **Show the result on the input image**





# Results





# Results





# Disadvantages

1. **Processing time**
2. **Ineffective results**